

✓ Task

Develop a comprehensive AI credit scoring and fraud detection system. This project will involve generating synthetic financial data, performing data preprocessing and feature engineering, conducting exploratory data analysis, training and evaluating multiple machine learning models (Logistic Regression, Random Forest, XGBoost), and implementing SHAP for model explainability. Additionally, the system will include a loan eligibility estimator, fraud detection logic, a user guidance engine for rejected applications, and a simulation of WhatsApp notifications for loan approval/rejection. The project will conclude with a summary highlighting its achievements and potential enhancements.

Project Introduction (Markdown)

Subtask:

Create markdown cells for the project title, problem description, solution overview, and business impact, as required for the Jupyter Notebook structure.

✓ AI Credit Scoring and Fraud Detection System

Problem Description

Traditional credit scoring methods often rely on limited historical financial data, which can lead to biased assessments, especially for individuals with thin credit files or those new to the credit system. This results in many creditworthy applicants being denied access to financial services, while others with high fraud risk may slip through. Existing fraud detection systems can also be reactive, identifying fraudulent activities only after they occur, leading to significant financial losses and reputational damage for financial institutions.

Key challenges include:

- **Inaccurate Credit Assessments:** Limited data points for certain demographics or new market entrants.
- **High Operational Costs:** Manual review processes for suspicious transactions are time-consuming and expensive.
- **Delayed Fraud Detection:** Reactive systems fail to prevent fraud in real-time, leading to irreversible losses.
- **Lack of Adaptability:** Static models struggle to keep pace with evolving fraud tactics and economic changes.
- **Customer Dissatisfaction:** Slow processing times and erroneous denials or approvals negatively impact customer experience.

Solution Overview

Our proposed AI Credit Scoring and Fraud Detection System leverages advanced machine learning techniques and alternative data sources to provide a more accurate, fair, and proactive approach to financial risk management. This system aims to overcome the limitations of traditional methods by:

Key Components:

- **Enhanced Credit Scoring Module:** Utilizes a wider array of data points, including transaction history, digital footprint, and behavioral patterns (with due privacy considerations), to generate a more comprehensive credit risk profile. This module will employ ensemble learning models to improve prediction accuracy and reduce bias.
- **Real-time Fraud Detection Engine:** Implements deep learning models, such as Recurrent Neural Networks (RNNs) and Convolutional Neural Networks (CNNs), to analyze transaction data in real-time. This engine will identify suspicious patterns and anomalies indicative of fraudulent activities, allowing for immediate intervention.
- **Adaptive Learning Capabilities:** The system will continuously learn from new data and feedback loops, allowing it to adapt to evolving fraud tactics and changes in economic conditions, ensuring the models remain robust and relevant.
- **Explainable AI (XAI) Framework:** Integrates XAI techniques to provide transparent and interpretable credit decisions and fraud alerts. This ensures compliance with regulatory requirements and builds trust with users and financial institutions.
- **User-friendly Interface:** Develops an intuitive dashboard for financial institutions to monitor credit scores, fraud alerts, and model performance, facilitating quick decision-making and operational efficiency.

Technological Stack:

- **Data Ingestion & Processing:** Apache Kafka, Apache Spark
- **Machine Learning Frameworks:** TensorFlow, PyTorch, Scikit-learn
- **Deployment:** Docker, Kubernetes
- **Database:** PostgreSQL, MongoDB

Business Impact

Implementing the AI Credit Scoring and Fraud Detection System is expected to deliver substantial business value and competitive advantages for financial institutions:

- **Reduced Financial Losses:** Proactive and real-time fraud detection will significantly minimize monetary losses due to fraudulent transactions.
- **Increased Revenue Streams:** More accurate credit scoring will enable financial institutions to approve a broader range of creditworthy applicants who might have been overlooked by traditional methods, expanding their customer base and loan portfolio.
- **Enhanced Operational Efficiency:** Automation of credit assessment and fraud detection processes will reduce manual workloads, cut operational costs, and free up human resources for more complex tasks.
- **Improved Customer Satisfaction:** Faster credit decisions and a reduction in erroneous denials will lead to a more positive customer experience, fostering loyalty and attracting new clients.
- **Stronger Compliance and Risk Management:** The Explainable AI (XAI) framework will provide transparency, aiding in compliance with regulatory requirements and strengthening the institution's overall risk management posture.
- **Competitive Advantage:** By leveraging cutting-edge AI, financial institutions can differentiate themselves in the market, offering superior services and a more secure environment for their customers.
- **Adaptability and Future-Proofing:** The adaptive learning capabilities ensure the system remains effective against evolving fraud patterns and market dynamics, protecting the institution's investments in the long term.

Library Setup

Subtask:

Install and import all necessary libraries: pandas, numpy, matplotlib, seaborn, scikit-learn, xgboost, and shap, ensuring the environment is ready for development.

Reasoning: First, I'll install all the necessary libraries using `pip`. I will install them in a single command for efficiency.

```
pip install pandas numpy matplotlib seaborn scikit-learn xgboost shap
```

```
Requirement already satisfied: pandas in /usr/local/lib/python3.12/dist-packages (2.2.2)
Requirement already satisfied: numpy in /usr/local/lib/python3.12/dist-packages (2.0.2)
Requirement already satisfied: matplotlib in /usr/local/lib/python3.12/dist-packages (3.10.0)
Requirement already satisfied: seaborn in /usr/local/lib/python3.12/dist-packages (0.13.2)
Requirement already satisfied: scikit-learn in /usr/local/lib/python3.12/dist-packages (1.6.1)
Requirement already satisfied: xgboost in /usr/local/lib/python3.12/dist-packages (3.1.3)
Requirement already satisfied: shap in /usr/local/lib/python3.12/dist-packages (0.50.0)
Requirement already satisfied: python-dateutil>=2.8.2 in /usr/local/lib/python3.12/dist-packages (from pandas) (2.9.0)
Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.12/dist-packages (from pandas) (2024.1)
Requirement already satisfied: tzdata>=2022.7 in /usr/local/lib/python3.12/dist-packages (from pandas) (2024.1)
Requirement already satisfied: contourpy>=1.0.1 in /usr/local/lib/python3.12/dist-packages (from matplotlib) (1.3.0)
Requirement already satisfied: cycler>=0.10 in /usr/local/lib/python3.12/dist-packages (from matplotlib) (0.12.1)
Requirement already satisfied: fonttools>=4.22.0 in /usr/local/lib/python3.12/dist-packages (from matplotlib) (4.53.0)
Requirement already satisfied: kiwisolver>=1.3.1 in /usr/local/lib/python3.12/dist-packages (from matplotlib) (1.4.5)
Requirement already satisfied: packaging>=20.0 in /usr/local/lib/python3.12/dist-packages (from matplotlib) (24.1)
Requirement already satisfied: pillow>=8 in /usr/local/lib/python3.12/dist-packages (from matplotlib) (10.4.0)
Requirement already satisfied: pyparsing>=2.3.1 in /usr/local/lib/python3.12/dist-packages (from matplotlib) (3.1.3)
Requirement already satisfied: scipy>=1.6.0 in /usr/local/lib/python3.12/dist-packages (from scikit-learn) (1.13.1)
Requirement already satisfied: joblib>=1.2.0 in /usr/local/lib/python3.12/dist-packages (from scikit-learn) (1.4.2)
Requirement already satisfied: threadpoolctl>=3.1.0 in /usr/local/lib/python3.12/dist-packages (from scikit-learn) (3.5.0)
Requirement already satisfied: nvidia-nccl-cu12 in /usr/local/lib/python3.12/dist-packages (from xgboost) (12.1.1)
Requirement already satisfied: tqdm>=4.27.0 in /usr/local/lib/python3.12/dist-packages (from shap) (4.66.1)
Requirement already satisfied: slicer==0.0.8 in /usr/local/lib/python3.12/dist-packages (from shap) (0.0.8)
Requirement already satisfied: numba>=0.54 in /usr/local/lib/python3.12/dist-packages (from shap) (0.60.0)
Requirement already satisfied: cloudpickle in /usr/local/lib/python3.12/dist-packages (from shap) (3.1.1)
Requirement already satisfied: typing-extensions in /usr/local/lib/python3.12/dist-packages (from shap) (4.11.0)
Requirement already satisfied: llvmlite<0.44,>=0.43.0dev0 in /usr/local/lib/python3.12/dist-packages (from numba) (0.43.0)
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.12/dist-packages (from python-dateutil) (1.17.0)
```

Reasoning: Now that the libraries are installed, I will import them as specified in the instructions to make them available for use in the notebook.

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import sklearn
import xgboost as xgb
import shap

print("All necessary libraries imported successfully.")

All necessary libraries imported successfully.
```

Synthetic Data Generation

Subtask:

Generate a realistic synthetic financial dataset for loan applications. This dataset will include features relevant for credit scoring and fraud detection, mimicking real-world financial data.

Reasoning: The first step is to generate the synthetic financial dataset as per the instructions. This involves defining the number of records, creating various financial and personal attributes using numpy, and populating a pandas DataFrame. I will start by generating the initial features and then introduce relationships and fraud patterns.

```
import pandas as pd
import numpy as np

# 1. Define the number of synthetic records to generate
n_records = 50000

# 3. Generate columns for various financial and personal attributes
np.random.seed(42) # For reproducibility

data = {
    'CustomerID': np.arange(1, n_records + 1),
    'Age': np.random.randint(18, 70, n_records), # Skewed slightly towards working adults
    'Income': np.random.lognormal(mean=10.5, sigma=0.8, size=n_records).round(2), # Realistic income distribution
    'LoanAmount': np.random.normal(loc=15000, scale=8000, size=n_records).round(2), # Realistic loan amount
    'CreditScore': np.random.randint(300, 850, n_records), # FICO score range
    'MonthsEmployed': np.random.randint(0, 300, n_records), # Up to 25 years
    'HasMortgage': np.random.choice([0, 1], size=n_records, p=[0.6, 0.4]), # 40% have a mortgage
    'HasDependents': np.random.choice([0, 1], size=n_records, p=[0.5, 0.5]), # 50% have dependents
    'LoanDuration': np.random.choice([12, 24, 36, 48, 60], size=n_records),
    'InterestRate': np.random.uniform(3.0, 18.0, n_records).round(2)
}

# Ensure LoanAmount and Income are non-negative
data['LoanAmount'] = np.maximum(data['LoanAmount'], 1000) # Minimum loan amount of 1000
data['Income'] = np.maximum(data['Income'], 15000) # Minimum income of 15000

# Create DataFrame
df = pd.DataFrame(data)

# 4. Introduce some relationships between variables for 'Default'
# Lower credit scores and lower income should correlate with a higher likelihood of Default.
# Base default rate
default_probability = 0.05 # 5% base default rate

# Adjust default probability based on CreditScore and Income
df['Default_Prob'] = default_probability
df['Default_Prob'] = df.apply(lambda row: row['Default_Prob'] + (850 - row['CreditScore']) / 850 * 0.15, axis=1)
df['Default_Prob'] = df.apply(lambda row: row['Default_Prob'] + (150000 - row['Income']) / 150000 * 0.1, axis=1)
df['Default_Prob'] = np.clip(df['Default_Prob'], 0.01, 0.9) # Clip probabilities to a realistic range

df['Default'] = (np.random.rand(n_records) < df['Default_Prob']).astype(int)

# Ensure a realistic imbalance for 'Default' (fewer defaults than non-defaults)
# If by chance too many defaults are generated, we can reduce some high-credit-score defaults.
while df['Default'].mean() > 0.15: # Aim for max 15% default rate
    high_cs_defaults = df[(df['Default'] == 1) & (df['CreditScore'] > 700)].index
    if len(high_cs_defaults) == 0: break
    to_change = np.random.choice(high_cs_defaults, size=min(int(len(high_cs_defaults)*0.1), 100), replace=True)
    df.loc[to_change, 'Default'] = 0

# 5. Introduce some patterns for 'IsFraud'
```

```
# Very high LoanAmount with low CreditScore could be a pattern.
# Base fraud rate
fraud_probability = 0.005 # 0.5% base fraud rate

# Adjust fraud probability based on LoanAmount and CreditScore
df['Fraud_Prob'] = fraud_probability
df['Fraud_Prob'] = df.apply(lambda row: row['Fraud_Prob'] + (row['LoanAmount'] / 50000 * 0.1) if row['L
df['Fraud_Prob'] = df.apply(lambda row: row['Fraud_Prob'] + (500 - row['CreditScore']) / 500 * 0.05 if
df['Fraud_Prob'] = np.clip(df['Fraud_Prob'], 0.001, 0.1) # Clip probabilities

df['IsFraud'] = (np.random.rand(n_records) < df['Fraud_Prob']).astype(int)

# Clean up temporary probability columns
df = df.drop(columns=['Default_Prob', 'Fraud_Prob'])

# 6. Display the first few rows of the generated DataFrame and its information
print("First 5 rows of the synthetic dataset:")
print(df.head())
print("\nDataFrame Information:")
df.info()
```

First 5 rows of the synthetic dataset:

	CustomerID	Age	Income	LoanAmount	CreditScore	MonthsEmployed	\
0	1	56	128734.18	17706.61	368	71	
1	2	69	27614.47	22220.38	330	166	
2	3	46	33019.96	27270.81	603	36	
3	4	32	15000.00	5118.90	479	217	
4	5	60	31098.53	18868.58	618	200	

	HasMortgage	HasDependents	LoanDuration	InterestRate	Default	IsFraud
0	1	1	60	13.63	0	0
1	0	0	60	16.31	0	0
2	1	0	24	4.83	0	0
3	0	1	60	11.93	0	0
4	0	0	24	7.24	0	0

DataFrame Information:

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 50000 entries, 0 to 49999
```

```
Data columns (total 12 columns):
```

#	Column	Non-Null Count	Dtype
0	CustomerID	50000 non-null	int64
1	Age	50000 non-null	int64
2	Income	50000 non-null	float64
3	LoanAmount	50000 non-null	float64
4	CreditScore	50000 non-null	int64
5	MonthsEmployed	50000 non-null	int64
6	HasMortgage	50000 non-null	int64
7	HasDependents	50000 non-null	int64
8	LoanDuration	50000 non-null	int64
9	InterestRate	50000 non-null	float64
10	Default	50000 non-null	int64
11	IsFraud	50000 non-null	int64

```
dtypes: float64(3), int64(9)
```

```
memory usage: 4.6 MB
```

✓ Data Preprocessing and Feature Engineering

Subtask:

Perform data cleaning, handle missing values, encode categorical variables, and create new features that are relevant for both credit risk assessment and fraud detection. This step also includes splitting the data into training and testing sets.

Reasoning: First, I will check for any missing values in the DataFrame (`df`). Then, I will create the new features as specified, handling potential division by zero for ratios, and finally, I will split the data into training and testing sets for both 'Default' and 'IsFraud' target variables.

```
print("Checking for missing values:")
print(df.isnull().sum())

# 2. Create new features
# Calculate DebtToIncomeRatio (LoanAmount / Income)
# Handle potential division by zero or very small incomes
df['DebtToIncomeRatio'] = df['LoanAmount'] / df['Income']
df['DebtToIncomeRatio'].replace([np.inf, -np.inf], np.nan, inplace=True)
median_debt_to_income = df['DebtToIncomeRatio'].median()
df['DebtToIncomeRatio'].fillna(median_debt_to_income, inplace=True)

# Calculate LoanToIncomeRatio
df['LoanToIncomeRatio'] = df['LoanAmount'] / df['Income']
df['LoanToIncomeRatio'].replace([np.inf, -np.inf], np.nan, inplace=True)
df['LoanToIncomeRatio'].fillna(median_debt_to_income, inplace=True) # Using same median for consistency

# Calculate CreditScorePerAge
df['CreditScorePerAge'] = df['CreditScore'] / df['Age']

# Approximate MonthlyLoanPayment (LoanAmount / LoanDuration)
df['MonthlyLoanPayment'] = df['LoanAmount'] / df['LoanDuration']

# Calculate LoanAmountPerDuration
df['LoanAmountPerDuration'] = df['LoanAmount'] / df['LoanDuration']

print("\nNew features created.")
print(df.head())

# 3. Define features (X) and target variables (y_default, y_fraud)
X = df.drop(columns=['CustomerID', 'Default', 'IsFraud'])
y_default = df['Default']
y_fraud = df['IsFraud']

print("\nFeatures (X) and target variables (y_default, y_fraud) defined.")
print(f"Shape of X: {X.shape}")
print(f"Shape of y_default: {y_default.shape}")
print(f"Shape of y_fraud: {y_fraud.shape}")

# 5. Split the data for default prediction
from sklearn.model_selection import train_test_split

X_train_default, X_test_default, y_train_default, y_test_default = train_test_split(X, y_default, test_size=0.2, random_state=42)

print("\nData split for default prediction:")
print(f"X_train_default shape: {X_train_default.shape}")
print(f"X_test_default shape: {X_test_default.shape}")
print(f"y_train_default shape: {y_train_default.shape}")
print(f"y_test_default shape: {y_test_default.shape}")

# 6. Split the data for fraud detection
X_train_fraud, X_test_fraud, y_train_fraud, y_test_fraud = train_test_split(X, y_fraud, test_size=0.2, random_state=42)

print("\nData split for fraud detection:")
print(f"X_train_fraud shape: {X_train_fraud.shape}")
print(f"X_test_fraud shape: {X_test_fraud.shape}")
```

```
print(f"y_train_fraud shape: {y_train_fraud.shape}")
print(f"y_test_fraud shape: {y_test_fraud.shape}")
```

0	0.137544	0.137544	6.571429
1	0.804664	0.804664	4.782609
2	0.825889	0.825889	13.108696
3	0.341260	0.341260	14.968750
4	0.606735	0.606735	10.300000

	MonthlyLoanPayment	LoanAmountPerDuration
0	295.110167	295.110167
1	370.339667	370.339667
2	1136.283750	1136.283750
3	85.315000	85.315000
4	786.190833	786.190833

Features (X) and target variables (y_default, y_fraud) defined.

Shape of X: (50000, 14)

Shape of y_default: (50000,)

Shape of y_fraud: (50000,)

Data split for default prediction:

X_train_default shape: (40000, 14)

X_test_default shape: (10000, 14)

y_train_default shape: (40000,)

y_test_default shape: (10000,)

Data split for fraud detection:

X_train_fraud shape: (40000, 14)

X_test_fraud shape: (10000, 14)

y_train_fraud shape: (40000,)

y_test_fraud shape: (10000,)

/tmp/ipython-input-1849722690.py:8: FutureWarning: A value is trying to be set on a copy of a DataFrame. The behavior will change in pandas 3.0. This inplace method will never work because the intermediate

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, in

```
df['DebtToIncomeRatio'].replace([np.inf, -np.inf], np.nan, inplace=True)
```

/tmp/ipython-input-1849722690.py:10: FutureWarning: A value is trying to be set on a copy of a DataFrame. The behavior will change in pandas 3.0. This inplace method will never work because the intermediate

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, in

```
df['DebtToIncomeRatio'].fillna(median_debt_to_income, inplace=True)
```

/tmp/ipython-input-1849722690.py:14: FutureWarning: A value is trying to be set on a copy of a DataFrame. The behavior will change in pandas 3.0. This inplace method will never work because the intermediate

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, in

```
df['LoanToIncomeRatio'].replace([np.inf, -np.inf], np.nan, inplace=True)
```

/tmp/ipython-input-1849722690.py:15: FutureWarning: A value is trying to be set on a copy of a DataFrame. The behavior will change in pandas 3.0. This inplace method will never work because the intermediate

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, in

```
df['LoanToIncomeRatio'].fillna(median_debt_to_income, inplace=True) # Using same median for consist
```

Reasoning: The previous code generated `FutureWarning` messages due to the use of `inplace=True` in pandas operations. To address these warnings and ensure compatibility with future pandas versions, I will modify the code to use direct assignment instead of inplace modifications for `replace` and `fillna` methods.

```
print("Checking for missing values:")
```

```
print(df.isnull().sum())

# 2. Create new features
# Calculate DebtToIncomeRatio (LoanAmount / Income)
# Handle potential division by zero or very small incomes
df['DebtToIncomeRatio'] = df['LoanAmount'] / df['Income']
df['DebtToIncomeRatio'] = df['DebtToIncomeRatio'].replace([np.inf, -np.inf], np.nan)
median_debt_to_income = df['DebtToIncomeRatio'].median()
df['DebtToIncomeRatio'] = df['DebtToIncomeRatio'].fillna(median_debt_to_income)

# Calculate LoanToIncomeRatio
df['LoanToIncomeRatio'] = df['LoanAmount'] / df['Income']
df['LoanToIncomeRatio'] = df['LoanToIncomeRatio'].replace([np.inf, -np.inf], np.nan)
df['LoanToIncomeRatio'] = df['LoanToIncomeRatio'].fillna(median_debt_to_income) # Using same median for

# Calculate CreditScorePerAge
df['CreditScorePerAge'] = df['CreditScore'] / df['Age']

# Approximate MonthlyLoanPayment (LoanAmount / LoanDuration)
df['MonthlyLoanPayment'] = df['LoanAmount'] / df['LoanDuration']

# Calculate LoanAmountPerDuration
df['LoanAmountPerDuration'] = df['LoanAmount'] / df['LoanDuration']

print("\nNew features created.")
print(df.head())

# 3. Define features (X) and target variables (y_default, y_fraud)
X = df.drop(columns=['CustomerID', 'Default', 'IsFraud'])
y_default = df['Default']
y_fraud = df['IsFraud']

print("\nFeatures (X) and target variables (y_default, y_fraud) defined.")
print(f"Shape of X: {X.shape}")
print(f"Shape of y_default: {y_default.shape}")
print(f"Shape of y_fraud: {y_fraud.shape}")

# 5. Split the data for default prediction
from sklearn.model_selection import train_test_split

X_train_default, X_test_default, y_train_default, y_test_default = train_test_split(X, y_default, test_

print("\nData split for default prediction:")
print(f"X_train_default shape: {X_train_default.shape}")
print(f"X_test_default shape: {X_test_default.shape}")
print(f"y_train_default shape: {y_train_default.shape}")
print(f"y_test_default shape: {y_test_default.shape}")

# 6. Split the data for fraud detection
X_train_fraud, X_test_fraud, y_train_fraud, y_test_fraud = train_test_split(X, y_fraud, test_size=0.2,

print("\nData split for fraud detection:")
print(f"X_train_fraud shape: {X_train_fraud.shape}")
print(f"X_test_fraud shape: {X_test_fraud.shape}")
print(f"y_train_fraud shape: {y_train_fraud.shape}")
print(f"y_test_fraud shape: {y_test_fraud.shape}")
```

```

LoanToIncomeRatio      0
CreditScorePerAge      0
MonthlyLoanPayment      0
LoanAmountPerDuration   0
dtype: int64

New features created.
CustomerID  Age  Income  LoanAmount  CreditScore  MonthsEmployed  \
0           1   56  128734.18    17706.61         368         71
1           2   69   27614.47    22220.38         330        166
2           3   46   33019.96    27270.81         603         36
3           4   32  15000.00     5118.90         479        217
4           5   60   31098.53    18868.58         618        200

HasMortgage  HasDependents  LoanDuration  InterestRate  Default  IsFraud  \
0            1              1           60         13.63         0         0
1            0              0           60         16.31         0         0
2            1              0           24          4.83         0         0
3            0              1           60         11.93         0         0
4            0              0           24          7.24         0         0

DebtToIncomeRatio  LoanToIncomeRatio  CreditScorePerAge  \
0          0.137544          0.137544          6.571429
1          0.804664          0.804664          4.782609
2          0.825889          0.825889         13.108696
3          0.341260          0.341260         14.968750
4          0.606735          0.606735         10.300000

MonthlyLoanPayment  LoanAmountPerDuration
0          295.110167          295.110167
1          370.339667          370.339667
2          1136.283750          1136.283750
3           85.315000           85.315000
4          786.190833          786.190833

Features (X) and target variables (y_default, y_fraud) defined.
Shape of X: (50000, 14)
Shape of y_default: (50000,)
Shape of y_fraud: (50000,)

Data split for default prediction:
X_train_default shape: (40000, 14)
X_test_default shape: (10000, 14)
y_train_default shape: (40000,)
y_test_default shape: (10000,)

Data split for fraud detection:
X_train_fraud shape: (40000, 14)
X_test_fraud shape: (10000, 14)
y_train_fraud shape: (40000,)
y_test_fraud shape: (10000,)

```

✓ Exploratory Data Analysis (EDA)

Subtask:

Conduct comprehensive EDA to understand data distributions, correlations, and potential insights. Include visualizations of key features and their relationships with the target variable (loan eligibility/default). Ensure all visualizations include clear legends.

Reasoning: I will start by generating histograms for all numerical features in the ☒ DataFrame to visualize their distributions, as requested in the first instruction for EDA. This will help in understanding the spread and central tendency of each feature.


```
print("Generating histograms for numerical features...")

# Identify numerical columns (all columns in X are numerical in this case)
numerical_cols = X.select_dtypes(include=np.number).columns

plt.figure(figsize=(20, 15))
for i, col in enumerate(numerical_cols):
    plt.subplot(4, 4, i + 1) # Adjust subplot grid as needed
    sns.histplot(X[col], kde=True)
    plt.title(f'Distribution of {col}')
    plt.xlabel(col)
    plt.ylabel('Frequency')
plt.tight_layout()
plt.show()

print("Histograms generated successfully.")
```

Generating histograms for numerical features...



Histograms generated successfully.

Reasoning: Now, I will generate a correlation heatmap to understand the relationships between all features in `X` and the target variables `y_default` and `y_fraud`, as specified in the second instruction.

```
print("Generating correlation heatmap...")


# Combine features and target variables into a single DataFrame for correlation calculation
df_corr = X.copy()
df_corr['Default'] = y_default
df_corr['IsFraud'] = y_fraud
```

```
# Calculate the correlation matrix
corr_matrix = df_corr.corr()

plt.figure(figsize=(16, 12))
sns.heatmap(corr_matrix, annot=True, cmap='coolwarm', fmt=".2f", linewidths=.5)
plt.title('Correlation Matrix of All Features and Target Variables')
plt.show()

print("Correlation heatmap generated successfully.")
```

Generating correlation heatmap...



Correlation heatmap generated successfully.

Reasoning: I will now generate box plots to visualize the relationship between the 'Default' target variable and key numerical features ('Age', 'Income', 'LoanAmount', 'CreditScore') as per the instructions, helping to understand how these features differ across default and non-default cases.

```
print("Generating box plots for 'Default' vs. key numerical features...")

# Create a temporary DataFrame for plotting 'Default' relationships
df_plot_default = X.copy()
df_plot_default['Default'] = y_default

key_numerical_features = ['Age', 'Income', 'LoanAmount', 'CreditScore']
```

```
plt.figure(figsize=(18, 12))
for i, feature in enumerate(key_numerical_features):
    plt.subplot(2, 2, i + 1)
    sns.boxplot(x='Default', y=feature, data=df_plot_default, palette='viridis')
    plt.title(f'{feature} vs. Default Status')
    plt.xlabel('Default Status (0=No, 1=Yes)')
    plt.ylabel(feature)
    plt.xticks([0, 1], ['No Default', 'Default'])
plt.tight_layout()
plt.show()

print("Box plots for 'Default' generated successfully.")
```

Generating box plots for 'Default' vs. key numerical features...

/tmp/ipython-input-2571994893.py:12: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x`

```
sns.boxplot(x='Default', y=feature, data=df_plot_default, palette='viridis')
```

/tmp/ipython-input-2571994893.py:12: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x`

```
sns.boxplot(x='Default', y=feature, data=df_plot_default, palette='viridis')
```

/tmp/ipython-input-2571994893.py:12: FutureWarning:


Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x`

```
sns.boxplot(x='Default', y=feature, data=df_plot_default, palette='viridis')
```

/tmp/ipython-input-2571994893.py:12: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x`

```
sns.boxplot(x='Default', y=feature, data=df_plot_default, palette='viridis')
```



Box plots for 'Default' generated successfully.

Reasoning: The previous code generated `FutureWarning` due to the use of `palette` without `hue` in `sns.boxplot`. I will modify the `sns.boxplot` call to include `hue='Default'` and `legend=False` to resolve this warning and align with future `seaborn` versions, while retaining the same visualization.

```
print("Generating box plots for 'Default' vs. key numerical features...")
```

```
# Create a temporary DataFrame for plotting 'Default' relationships
```

```
df_plot_default = X.copy()
```

```
df_plot_default['Default'] = y_default
```

```


key_numerical_features = ['Age', 'Income', 'LoanAmount', 'CreditScore']

plt.figure(figsize=(18, 12))
for i, feature in enumerate(key_numerical_features):
    plt.subplot(2, 2, i + 1)
    sns.boxplot(x='Default', y=feature, data=df_plot_default, hue='Default', palette='viridis', legend=
plt.title(f'{feature} vs. Default Status')
plt.xlabel('Default Status (0=No, 1=Yes)')
plt.ylabel(feature)
plt.xticks([0, 1], ['No Default', 'Default'])
plt.tight_layout()
plt.show()

print("Box plots for 'Default' generated successfully.")

```

Generating box plots for 'Default' vs. key numerical features...



Box plots for 'Default' generated successfully.

Reasoning: Now, I will generate box plots to visualize the relationship between the 'IsFraud' target variable and key numerical features ('Age', 'Income', 'LoanAmount', 'CreditScore') as specified in the fourth instruction, helping to understand how these features differ across fraudulent and non-fraudulent cases.

```

print("Generating box plots for 'IsFraud' vs. key numerical features...")

# Create a temporary DataFrame for plotting 'IsFraud' relationships
df_plot_fraud = X.copy()
df_plot_fraud['IsFraud'] = y_fraud

key_numerical_features = ['Age', 'Income', 'LoanAmount', 'CreditScore']


plt.figure(figsize=(18, 12))
for i, feature in enumerate(key_numerical_features):

```

```
plt.subplot(2, 2, i + 1)
sns.boxplot(x='IsFraud', y=feature, data=df_plot_fraud, hue='IsFraud', palette='magma', legend=False)
plt.title(f'{feature} vs. Fraud Status')
plt.xlabel('Fraud Status (0=No, 1=Yes)')
plt.ylabel(feature)
plt.xticks([0, 1], ['No Fraud', 'Fraud'])
plt.tight_layout()
plt.show()

print("Box plots for 'IsFraud' generated successfully.")
```

Generating box plots for 'IsFraud' vs. key numerical features...



Box plots for 'IsFraud' generated successfully.

Reasoning: Next, I will analyze the distribution of categorical features ('HasMortgage', 'HasDependents', 'LoanDuration') with respect to 'Default' and 'IsFraud' using count plots, as specified in the fifth instruction of the EDA subtask. This will help understand their influence on the target variables.

```
print("Generating count plots for categorical features vs. Default and IsFraud...")

categorical_features = ['HasMortgage', 'HasDependents', 'LoanDuration']


# Plotting for 'Default'
plt.figure(figsize=(18, 6))
for i, feature in enumerate(categorical_features):
    plt.subplot(1, 3, i + 1)
    sns.countplot(x=feature, hue='Default', data=df, palette='viridis')
    plt.title(f'{feature} vs. Default Status')
    plt.xlabel(feature)
    plt.ylabel('Count')
    plt.legend(title='Default', labels=['No Default', 'Default'])
plt.tight_layout()
```

```
plt.show()

# Plotting for 'IsFraud'
plt.figure(figsize=(18, 6))
for i, feature in enumerate(categorical_features):
    plt.subplot(1, 3, i + 1)
    sns.countplot(x=feature, hue='IsFraud', data=df, palette='magma')
    plt.title(f'{feature} vs. Fraud Status')
    plt.xlabel(feature)
    plt.ylabel('Count')
    plt.legend(title='IsFraud', labels=['No Fraud', 'Fraud'])
plt.tight_layout()
plt.show()

print("Count plots for categorical features generated successfully.")
```

Generating count plots for categorical features vs. Default and IsFraud...



Count plots for categorical features generated successfully.

✓ Model Training (Logistic Regression, Random Forest, XGBoost)

Subtask:

Train a Logistic Regression model as a baseline, a Random Forest classifier, and an XGBoost classifier for credit risk scoring. These models will predict the likelihood of loan default.

Reasoning: I will import the required machine learning models from their respective libraries, then instantiate and train each model (Logistic Regression, Random Forest, and XGBoost) using the preprocessed training data for default prediction. Finally, I will print a confirmation message.

```

from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier
from xgboost import XGBClassifier

print("Training Logistic Regression model...")
lr_model = LogisticRegression(random_state=42, solver='liblinear') # Added solver for better converger
lr_model.fit(X_train_default, y_train_default)
print("Logistic Regression model trained.")

print("\nTraining Random Forest Classifier model...")
rf_model = RandomForestClassifier(random_state=42, n_jobs=-1) # n_jobs=-1 uses all available cores
rf_model.fit(X_train_default, y_train_default)
print("Random Forest Classifier model trained.")

print("\nTraining XGBoost Classifier model...")
xgb_model = XGBClassifier(random_state=42, use_label_encoder=False, eval_metric='logloss') # use_label
xgb_model.fit(X_train_default, y_train_default)
print("XGBoost Classifier model trained.")

print("\nAll credit risk prediction models (Logistic Regression, Random Forest, XGBoost) have been tra

```

```

Training Logistic Regression model...
Logistic Regression model trained.

Training Random Forest Classifier model...
Random Forest Classifier model trained.

Training XGBoost Classifier model...
/usr/local/lib/python3.12/dist-packages/xgboost/training.py:199: UserWarning: [15:10:58] WARNING: /work
Parameters: { "use_label_encoder" } are not used.

    bst.update(dtrain, iteration=i, fobj=obj)
XGBoost Classifier model trained.

All credit risk prediction models (Logistic Regression, Random Forest, XGBoost) have been trained succe

```

Reasoning: The previous code block produced a UserWarning regarding the `use_label_encoder` parameter in `XGBClassifier`, indicating it's no longer used. I will remove this parameter from the `XGBClassifier` instantiation to prevent the warning.

```

from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier
from xgboost import XGBClassifier

print("Training Logistic Regression model...")
lr_model = LogisticRegression(random_state=42, solver='liblinear') # Added solver for better convergenc
lr_model.fit(X_train_default, y_train_default)
print("Logistic Regression model trained.")

print("\nTraining Random Forest Classifier model...")
rf_model = RandomForestClassifier(random_state=42, n_jobs=-1) # n_jobs=-1 uses all available cores
rf_model.fit(X_train_default, y_train_default)
print("Random Forest Classifier model trained.")

print("\nTraining XGBoost Classifier model...")
xgb_model = XGBClassifier(random_state=42, eval_metric='logloss') # Removed use_label_encoder=False as
xgb_model.fit(X_train_default, y_train_default)
print("XGBoost Classifier model trained.")

print("\nAll credit risk prediction models (Logistic Regression, Random Forest, XGBoost) have been trai

```

```
Training Logistic Regression model...
Logistic Regression model trained.
```

```
Training Random Forest Classifier model...
Random Forest Classifier model trained.
```

```
Training XGBoost Classifier model...
XGBoost Classifier model trained.
```

All credit risk prediction models (Logistic Regression, Random Forest, XGBoost) have been trained successfully.

✓ Model Evaluation

Subtask:

Evaluate the performance of all trained models using ROC curves, AUC scores, Confusion Matrices, Precision, Recall, and F1-scores. Visualize ROC curves for comparison across models. Ensure all visualizations include clear legends.

Reasoning: I will start by importing all the necessary evaluation metrics and plotting functions from

`sklearn.metrics` to prepare for model evaluation.

```
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score, roc_curve, auc, co
import matplotlib.pyplot as plt
import seaborn as sns
```

```
print("Necessary evaluation metrics and plotting functions imported.")
```

```
Necessary evaluation metrics and plotting functions imported.
```

Reasoning: Now, I will evaluate the Logistic Regression model for default prediction by making predictions on the test set, calculating accuracy, precision, recall, F1-score, and AUC, and then visualizing its confusion matrix.

```
print("Evaluating Logistic Regression Model...")

y_pred_lr = lr_model.predict(X_test_default)
y_proba_lr = lr_model.predict_proba(X_test_default)[:, 1] # Probability of positive class (default)

# Calculate metrics
accuracy_lr = accuracy_score(y_test_default, y_pred_lr)
precision_lr = precision_score(y_test_default, y_pred_lr)
recall_lr = recall_score(y_test_default, y_pred_lr)
f1_lr = f1_score(y_test_default, y_pred_lr)
auc_lr = auc(roc_curve(y_test_default, y_proba_lr)[0], roc_curve(y_test_default, y_proba_lr)[1])

print(f"\nLogistic Regression Metrics:")
print(f"Accuracy: {accuracy_lr:.4f}")
print(f"Precision: {precision_lr:.4f}")
print(f"Recall: {recall_lr:.4f}")
print(f"F1-Score: {f1_lr:.4f}")
print(f"AUC: {auc_lr:.4f}")

# Confusion Matrix
cm_lr = confusion_matrix(y_test_default, y_pred_lr)
plt.figure(figsize=(6, 5))
sns.heatmap(cm_lr, annot=True, fmt='d', cmap='Blues', cbar=False,
            xticklabels=['No Default', 'Default'], yticklabels=['No Default', 'Default'])
```

```
plt.title('Confusion Matrix - Logistic Regression')
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.show()
```

```
print("Logistic Regression evaluation complete.")
```

Evaluating Logistic Regression Model...

Logistic Regression Metrics:

Accuracy: 0.8549


Precision: 0.0000

Recall: 0.0000

F1-Score: 0.0000

AUC: 0.6260

/usr/local/lib/python3.12/dist-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Precision, Recall and F1-Score are ill-defined with zero true or predicted counts. _warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))



Logistic Regression evaluation complete.

Reasoning: The previous code block produced an `UndefinedMetricWarning` because the Logistic Regression model predicted no positive samples, leading to zero precision, recall, and F1-score. To resolve this warning and ensure robustness, I will explicitly set the `zero_division` parameter to 0 in the `precision_score`, `recall_score`, and `f1_score` functions. This will ensure that when no positive predictions are made, the score remains 0.0 without raising a warning.

```
print("Evaluating Logistic Regression Model...")

y_pred_lr = lr_model.predict(X_test_default)
y_proba_lr = lr_model.predict_proba(X_test_default)[:, 1] # Probability of positive class (default)

# Calculate metrics
accuracy_lr = accuracy_score(y_test_default, y_pred_lr)
precision_lr = precision_score(y_test_default, y_pred_lr, zero_division=0)
recall_lr = recall_score(y_test_default, y_pred_lr, zero_division=0)
f1_lr = f1_score(y_test_default, y_pred_lr, zero_division=0)
auc_lr = auc(roc_curve(y_test_default, y_proba_lr)[0], roc_curve(y_test_default, y_proba_lr)[1])
```

```

print(f"\nLogistic Regression Metrics:")
print(f"Accuracy: {accuracy_lr:.4f}")
print(f"Precision: {precision_lr:.4f}")
print(f"Recall: {recall_lr:.4f}")
print(f"F1-Score: {f1_lr:.4f}")
print(f"AUC: {auc_lr:.4f}")


# Confusion Matrix
cm_lr = confusion_matrix(y_test_default, y_pred_lr)
plt.figure(figsize=(6, 5))
sns.heatmap(cm_lr, annot=True, fmt='d', cmap='Blues', cbar=False,
            xticklabels=['No Default', 'Default'], yticklabels=['No Default', 'Default'])
plt.title('Confusion Matrix - Logistic Regression')
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.show()

print("Logistic Regression evaluation complete.")

```

Evaluating Logistic Regression Model...

Logistic Regression Metrics:
Accuracy: 0.8549
Precision: 0.0000
Recall: 0.0000
F1-Score: 0.0000
AUC: 0.6260



Logistic Regression evaluation complete.

Reasoning: Now, I will evaluate the Random Forest Classifier model for default prediction by making predictions on the test set, calculating accuracy, precision, recall, F1-score, and AUC, and then visualizing its confusion matrix, similar to the Logistic Regression model evaluation.

```

print("Evaluating Random Forest Classifier Model...")

y_pred_rf = rf_model.predict(X_test_default)

```

```

y_proba_rf = rf_model.predict_proba(X_test_default)[: , 1] # Probability of positive class (default)

# Calculate metrics
accuracy_rf = accuracy_score(y_test_default, y_pred_rf)
precision_rf = precision_score(y_test_default, y_pred_rf, zero_division=0)
recall_rf = recall_score(y_test_default, y_pred_rf, zero_division=0)
f1_rf = f1_score(y_test_default, y_pred_rf, zero_division=0)
auc_rf = auc(roc_curve(y_test_default, y_proba_rf)[0], roc_curve(y_test_default, y_proba_rf)[1])

print(f"\nRandom Forest Classifier Metrics:")
print(f"Accuracy: {accuracy_rf:.4f}")
print(f"Precision: {precision_rf:.4f}")
print(f"Recall: {recall_rf:.4f}")
print(f"F1-Score: {f1_rf:.4f}")
print(f"AUC: {auc_rf:.4f}")

# Confusion Matrix
cm_rf = confusion_matrix(y_test_default, y_pred_rf)
plt.figure(figsize=(6, 5))
sns.heatmap(cm_rf, annot=True, fmt='d', cmap='Greens', cbar=False,
            xticklabels=['No Default', 'Default'], yticklabels=['No Default', 'Default'])
plt.title('Confusion Matrix - Random Forest Classifier')
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.show()

print("Random Forest Classifier evaluation complete.")

```

Evaluating Random Forest Classifier Model...

Random Forest Classifier Metrics:


Accuracy: 0.8547

Precision: 0.0000

Recall: 0.0000

F1-Score: 0.0000

AUC: 0.5901



Random Forest Classifier evaluation complete.

Reasoning: Next, I will evaluate the XGBoost Classifier model for default prediction by making predictions on the test set, calculating accuracy, precision, recall, F1-score, and AUC, and then visualizing its confusion matrix, following the pattern of the previous model evaluations.

```
print("Evaluating XGBoost Classifier Model...")

y_pred_xgb = xgb_model.predict(X_test_default)
y_proba_xgb = xgb_model.predict_proba(X_test_default)[:, 1] # Probability of positive class (default)

# Calculate metrics
accuracy_xgb = accuracy_score(y_test_default, y_pred_xgb)
precision_xgb = precision_score(y_test_default, y_pred_xgb, zero_division=0)
recall_xgb = recall_score(y_test_default, y_pred_xgb, zero_division=0)
f1_xgb = f1_score(y_test_default, y_pred_xgb, zero_division=0)
auc_xgb = auc(roc_curve(y_test_default, y_proba_xgb)[0], roc_curve(y_test_default, y_proba_xgb)[1])

print(f"\nXGBoost Classifier Metrics:")
print(f"Accuracy: {accuracy_xgb:.4f}")
print(f"Precision: {precision_xgb:.4f}")
print(f"Recall: {recall_xgb:.4f}")
print(f"F1-Score: {f1_xgb:.4f}")
print(f"AUC: {auc_xgb:.4f}")

# Confusion Matrix
cm_xgb = confusion_matrix(y_test_default, y_pred_xgb)
plt.figure(figsize=(6, 5))
sns.heatmap(cm_xgb, annot=True, fmt='d', cmap='Blues', cbar=False,
            xticklabels=['No Default', 'Default'], yticklabels=['No Default', 'Default'])
plt.title('Confusion Matrix - XGBoost Classifier')
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.show()

print("XGBoost Classifier evaluation complete.")
```

Evaluating XGBoost Classifier Model...

XGBoost Classifier Metrics:


Accuracy: 0.8525

Precision: 0.2000

Recall: 0.0055

F1-Score: 0.0107

AUC: 0.5921



XGBoost Classifier evaluation complete.

Reasoning: I will now plot the ROC curves for all three models (Logistic Regression, Random Forest, and XGBoost) on a single graph, including a diagonal dashed line for reference, and ensure all labels and a legend are present for clear comparison.

```
print("Plotting ROC curves for all models...")

plt.figure(figsize=(10, 8))

# Plot ROC for Logistic Regression
fpr_lr, tpr_lr, _ = roc_curve(y_test_default, y_proba_lr)
plt.plot(fpr_lr, tpr_lr, label=f'Logistic Regression (AUC = {auc_lr:.4f})', color='blue')

# Plot ROC for Random Forest
fpr_rf, tpr_rf, _ = roc_curve(y_test_default, y_proba_rf)
plt.plot(fpr_rf, tpr_rf, label=f'Random Forest (AUC = {auc_rf:.4f})', color='green')

# Plot ROC for XGBoost
fpr_xgb, tpr_xgb, _ = roc_curve(y_test_default, y_proba_xgb)
plt.plot(fpr_xgb, tpr_xgb, label=f'XGBoost (AUC = {auc_xgb:.4f})', color='red')


# Plot diagonal line for random classifier
plt.plot([0, 1], [0, 1], 'k--', label='Random Classifier')

plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curve Comparison for Default Prediction')
plt.legend()
```

```
plt.grid(True)
plt.show()
```

```
print("ROC curves plotted successfully.")
```

Plotting ROC curves for all models...



ROC curves plotted successfully.

✓ SHAP Explainability

Subtask:

Implement SHAP (SHapley Additive exPlanations) for the best-performing model (Logistic Regression). Generate global feature importance plots (e.g., summary plot) and individual user explainability plots (e.g., force plot) for specific loan applications to understand model predictions. Ensure all visualizations include clear legends.

Reasoning: I will initialize a SHAP LinearExplainer for the Logistic Regression model, using the training data `X_train_default` to determine feature distributions, as instructed. This is appropriate for linear models like Logistic Regression.

```
print("Initializing SHAP Explainer for Logistic Regression...")
```

```
# 1. Initialize a SHAP Explainer for the lr_model
# For linear models like Logistic Regression, LinearExplainer is suitable.
explainer = shap.LinearExplainer(lr_model, X_train_default, feature_names=X_train_default.columns)

print("SHAP Explainer initialized successfully.")
```

```
Initializing SHAP Explainer for Logistic Regression...
SHAP Explainer initialized successfully.
```

Reasoning: Now that the SHAP explainer is initialized, I will calculate the SHAP values for the `X_test_default` dataset using the explainer, which will be used for subsequent plotting.

```
print("Calculating SHAP values for the test dataset...")

# 2. Calculate the SHAP values for the X_test_default dataset
shap_values = explainer.shap_values(X_test_default)

print("SHAP values calculated successfully.")
```

```
Calculating SHAP values for the test dataset...
SHAP values calculated successfully.
```


Reasoning: I will now generate a SHAP summary plot as a bar chart to visualize the global feature importance of the Logistic Regression model, as specified in the instructions. This will provide an overview of which features contribute most to the model's predictions.

```
print("Generating SHAP summary plot (bar) for global feature importance...")

# 3. Generate a SHAP summary plot (bar) for global feature importance
shap.summary_plot(shap_values, X_test_default, plot_type="bar", show=False)
plt.title('SHAP Global Feature Importance (Bar Plot)')
plt.tight_layout()
plt.show()

print("SHAP summary plot (bar) generated successfully.")
```

Generating SHAP summary plot (bar) for global feature importance...
 /tmp/ipython-input-2526989635.py:4: FutureWarning: The NumPy global RNG was seeded by calling `np.random.seed`
 shap.summary_plot(shap_values, X_test_default, plot_type="bar", show=False)



SHAP summary plot (bar) generated successfully.

Reasoning: The previous code produced a `FutureWarning` related to `np.random.seed` originating from within the SHAP library's internal operations. Since this warning does not halt execution and is not directly addressable by modifying the `shap.summary_plot` call, I will suppress this specific warning using `warnings.filterwarnings` to clean up the output as requested.

```
import warnings # Import warnings module


print("Generating SHAP summary plot (bar) for global feature importance...")

# Temporarily ignore the FutureWarning related to numpy.random.seed
with warnings.catch_warnings():
    warnings.simplefilter("ignore", FutureWarning)

    # 3. Generate a SHAP summary plot (bar) for global feature importance
    shap.summary_plot(shap_values, X_test_default, plot_type="bar", show=False)
    plt.title('SHAP Global Feature Importance (Bar Plot)')
    plt.tight_layout()
    plt.show()
```

```
print("SHAP summary plot (bar) generated successfully.")
```

Generating SHAP summary plot (bar) for global feature importance...



SHAP summary plot (bar) generated successfully.

Reasoning: Following the instructions, I will now generate a SHAP summary plot as a beeswarm plot to visualize feature impact and direction across the dataset for the Logistic Regression model. This will provide a more detailed view of how each feature influences individual predictions.

```
import warnings


print("Generating SHAP summary plot (beeswarm) for feature impact and direction...")

# Temporarily ignore the FutureWarning related to numpy.random.seed
with warnings.catch_warnings():
    warnings.simplefilter("ignore", FutureWarning)

    # 4. Generate another SHAP summary plot, this time as a beeswarm plot
    shap.summary_plot(shap_values, X_test_default, plot_type="beeswarm", show=False)
plt.title('SHAP Feature Impact and Direction (Beeswarm Plot)')
plt.tight_layout()
plt.show()

print("SHAP summary plot (beeswarm) generated successfully.")
```

Generating SHAP summary plot (beeswarm) for feature impact and direction...



SHAP summary plot (beeswarm) generated successfully.

Reasoning: I will now select a specific instance from the `X_test_default` dataset and generate a SHAP force plot for it, as required by the instructions. This will provide an individual explanation for the model's prediction on that particular loan application.

```
print("Generating SHAP force plot for a specific instance...")
```

```
# 5. Select a specific instance for individual explanation (e.g., the first instance)
instance_to_explain_idx = 0
instance_to_explain = X_test_default.iloc[instance_to_explain_idx]
shap_values_instance = shap_values[instance_to_explain_idx]

# 6. Generate a SHAP force plot for the selected instance
# shap.initjs() is usually called once in the notebook setup for JS visualizations
shap.initjs()
print(f"Force plot for instance {instance_to_explain_idx}:")
shap.force_plot(explainer.expected_value, shap_values_instance, instance_to_explain)

print("SHAP force plot generated successfully.")
```

Generating SHAP force plot for a specific instance...



Force plot for instance 0:
SHAP force plot generated successfully.

✓ Loan Eligibility Estimator

Subtask:

Develop a function or class that takes a user's financial profile, predicts their credit risk score, determines loan eligibility based on a defined threshold, calculates the eligible loan amount, and computes EMI and tenure logic based on the approved amount and risk.

Reasoning: I will define the `estimate_loan_eligibility` function, incorporating all the specified logic for predicting default probability, determining eligibility, calculating loan amount, and computing EMI and tenure. Finally, I will test this function with an example user profile to demonstrate its functionality and display the results.

```
def estimate_loan_eligibility(user_profile, model, eligibility_threshold=0.25, min_loan_amount=1000, n
    """
    Estimates loan eligibility, calculates eligible loan amount, and computes EMI and tenure.

    Args:
        user_profile (pd.Series): A Series representing the new user's financial profile.
        model: The trained machine learning model (e.g., Logistic Regression) for default prediction.
        eligibility_threshold (float): Maximum default probability for a loan to be eligible.
        min_loan_amount (int): Minimum loan amount to be considered.
        max_loan_amount (int): Maximum loan amount that can be approved.

    Returns:
        dict: A dictionary containing eligibility status, predicted default probability,
              eligible loan amount, monthly EMI, and loan tenure.
    """
    # Ensure user_profile is in the correct format for prediction (2D array/DataFrame)
    user_profile_df = pd.DataFrame([user_profile], columns=user_profile.index)

    # Predict the probability of default
    default_probability = model.predict_proba(user_profile_df)[0, 1][0]

    # Determine loan eligibility
    is_eligible = default_probability < eligibility_threshold

    eligible_loan_amount = 0
    monthly_emi = 0
    loan_tenure = 0
```

```

if is_eligible:
    # Calculate eligible loan amount
    # Example logic: Income * (CreditScore / 850) * a_factor
    # The factor can be tuned. Let's use 0.5 as a starting point.
    a_factor = 0.5
    calculated_loan_amount = user_profile['Income'] * (user_profile['CreditScore'] / 850) * a_factor
    eligible_loan_amount = np.clip(calculated_loan_amount, min_loan_amount, max_loan_amount)

    # Determine loan tenure (in months)
    # Use existing LoanDuration or a range based on risk
    # For simplicity, let's use the user's preferred LoanDuration for now
    # Could be adjusted based on default_probability: higher risk -> shorter tenure
    loan_tenure = int(user_profile['LoanDuration'])
    n_payments = loan_tenure # Number of monthly payments

    # Calculate monthly interest rate from annual percentage rate
    annual_interest_rate_percent = user_profile['InterestRate']
    monthly_interest_rate = (annual_interest_rate_percent / 100) / 12

    # Compute Monthly EMI using the formula:  $P * r * (1 + r)^n / ((1 + r)^n - 1)$ 
    principal = eligible_loan_amount
    r = monthly_interest_rate
    n = n_payments

    if r > 0:
        monthly_emi = principal * r * (1 + r)**n / (((1 + r)**n) - 1)
    else: # Handle zero interest rate
        monthly_emi = principal / n

else:
    eligible_loan_amount = 0
    monthly_emi = 0
    loan_tenure = 0

return {
    'is_eligible': bool(is_eligible),
    'default_probability': float(default_probability),
    'eligible_loan_amount': float(eligible_loan_amount),
    'monthly_emi': float(monthly_emi),
    'loan_tenure_months': int(loan_tenure)
}

# Test the function with an example user profile
print("\nTesting the loan eligibility estimator with an example user profile...")
example_user_profile = X_test_default.iloc[0]

# Using the best performing model based on AUC, which was Logistic Regression (lr_model)
loan_status = estimate_loan_eligibility(example_user_profile, lr_model)

print("\nLoan Eligibility Estimation for example user:")
for key, value in loan_status.items():
    print(f"  {key}: {value}")

```

Testing the loan eligibility estimator with an example user profile...

```

Loan Eligibility Estimation for example user:
is_eligible: True
default_probability: 0.23497070595827022
eligible_loan_amount: 7609.949788235294
monthly_emi: 258.46673389781216
loan_tenure_months: 36

```

✓ Fraud Detection Logic

Subtask:

Implement logic to detect abnormal transaction patterns and risky user financial behavior, identifying outliers or suspicious activities within the synthetic dataset to flag potential fraud users.

Reasoning: I will train an XGBoost Classifier for fraud detection using the provided training data (`X_train_fraud`, `y_train_fraud`) and store the trained model in the `fraud_model` variable, as per the first instruction.

```
from xgboost import XGBClassifier

print("Training XGBoost Classifier for fraud detection...")
# 1. Train an XGBoost Classifier for fraud detection
fraud_model = XGBClassifier(random_state=42, eval_metric='logloss') # Removed use_label_encoder=False
fraud_model.fit(X_train_fraud, y_train_fraud)
print("XGBoost Classifier for fraud detection trained successfully.")
```

```
Training XGBoost Classifier for fraud detection...
XGBoost Classifier for fraud detection trained successfully.
```

Reasoning: Now that the fraud detection model is trained, I will use it to predict the probabilities and class labels for fraud on the `X_test_fraud` dataset, as requested by the second instruction.

```
print("Predicting fraud probabilities and labels on the test dataset...")

y_proba_fraud = fraud_model.predict_proba(X_test_fraud)[:, 1] # Probability of fraud
y_pred_fraud = fraud_model.predict(X_test_fraud) # Predicted fraud labels

print("Fraud probabilities and labels predicted successfully.")
```

```
Predicting fraud probabilities and labels on the test dataset...
Fraud probabilities and labels predicted successfully.
```

Reasoning: I will calculate and print the accuracy, precision, recall, F1-score, and AUC for the fraud detection model using `y_test_fraud`, `y_pred_fraud`, and `y_proba_fraud`, ensuring to use `zero_division=0` for precision, recall, and F1-score as specified in the third instruction.

```
print("Evaluating Fraud Detection Model...")

# Calculate metrics for fraud detection
accuracy_fraud = accuracy_score(y_test_fraud, y_pred_fraud)
precision_fraud = precision_score(y_test_fraud, y_pred_fraud, zero_division=0)
recall_fraud = recall_score(y_test_fraud, y_pred_fraud, zero_division=0)
f1_fraud = f1_score(y_test_fraud, y_pred_fraud, zero_division=0)
auc_fraud = auc(roc_curve(y_test_fraud, y_proba_fraud)[0], roc_curve(y_test_fraud, y_proba_fraud)[1])

print(f"\nFraud Detection Metrics:")
print(f"Accuracy: {accuracy_fraud:.4f}")
print(f"Precision: {precision_fraud:.4f}")
print(f"Recall: {recall_fraud:.4f}")
print(f"F1-Score: {f1_fraud:.4f}")
print(f"AUC: {auc_fraud:.4f}")

print("Fraud detection model evaluation metrics calculated successfully.")
```

Evaluating Fraud Detection Model...

Fraud Detection Metrics:
Accuracy: 0.9900
Precision: 1.0000
Recall: 0.0099
F1-Score: 0.0196
AUC: 0.6341
Fraud detection model evaluation metrics calculated successfully.


Reasoning: I will generate and display a confusion matrix heatmap for the fraud detection model using `y_test_fraud` and `y_pred_fraud`, adding clear labels and tick labels as specified in the fourth instruction.

```
print("Generating Confusion Matrix for Fraud Detection Model...")

# Confusion Matrix for fraud detection
cm_fraud = confusion_matrix(y_test_fraud, y_pred_fraud)
plt.figure(figsize=(6, 5))
sns.heatmap(cm_fraud, annot=True, fmt='d', cmap='Reds', cbar=False,
            xticklabels=['No Fraud', 'Fraud'], yticklabels=['No Fraud', 'Fraud'])
plt.title('Confusion Matrix - Fraud Detection Model')
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.show()

print("Confusion Matrix for Fraud Detection Model generated successfully.")
```

Generating Confusion Matrix for Fraud Detection Model...



Confusion Matrix for Fraud Detection Model generated successfully.

✓ User Guidance Engine

Subtask:

Create a system to provide actionable suggestions and improvement tips for users who are rejected for a loan. This guidance will be based on the model's predictions and SHAP explanations, helping users understand how to improve their eligibility.

Reasoning: I will define the `generate_guidance` function, which will analyze a user's default probability and SHAP values to provide personalized suggestions for improving loan eligibility for rejected applications. I will then test this function with an example user likely to be rejected and display the generated guidance.

```
def generate_guidance(user_profile, default_probability, shap_values, feature_names, eligibility_thres
    """
    Generates actionable suggestions for users with rejected loan applications.

    Args:
        user_profile (pd.Series): The user's financial profile.
        default_probability (float): The predicted probability of default.
        shap_values (np.ndarray): SHAP values for the user's profile.
        feature_names (list): List of feature names corresponding to SHAP values.
        eligibility_threshold (float): The threshold above which an application is rejected.

    Returns:
        list: A list of strings containing personalized guidance.
    """
    guidance = []

    if default_probability >= eligibility_threshold:
        guidance.append(f"Your loan application was rejected due to a high estimated default probability. Consider improving your credit score and reducing debt-to-income ratio.")
```

```

guidance.append("Here are some personalized suggestions to improve your loan eligibility:")

# Combine SHAP values with feature names for easier analysis
shap_impact = pd.Series(shap_values, index=feature_names)

# Identify features with the most significant negative impact (large positive SHAP values for
# Sorting by SHAP values in descending order to prioritize features that push towards default
negative_impact_features = shap_impact.sort_values(ascending=False)

# Provide specific guidance based on the top negatively impacting features
for feature, impact in negative_impact_features.items():
    if impact > 0.01: # Consider features with a notable positive SHAP value
        if feature == 'CreditScore' and user_profile[feature] < 650:
            guidance.append(f" - Your Credit Score ({user_profile[feature]:.0f}) is a major fa
        elif feature == 'DebtToIncomeRatio' and user_profile[feature] > 0.4:
            guidance.append(f" - Your Debt-to-Income Ratio ({user_profile[feature]:.2f}) is hi
        elif feature == 'LoanAmount' and user_profile[feature] > 30000:
            guidance.append(f" - The requested Loan Amount ({user_profile[feature]:.2f}) might
        elif feature == 'Income' and user_profile[feature] < 20000:
            guidance.append(f" - Your Income ({user_profile[feature]:.2f}) is impacting your
        elif feature == 'InterestRate' and user_profile[feature] > 10:
            guidance.append(f" - The high Interest Rate ({user_profile[feature]:.2f}%) on your
        elif feature == 'MonthsEmployed' and user_profile[feature] < 12:
            guidance.append(f" - Your Months Employed ({user_profile[feature]:.0f}) indicates
        # Add more rules for other features as needed

    if len(guidance) == 2: # Only generic rejection message and "Here are..." are present
        guidance.append(" - We recommend reviewing your overall financial health to identify areas

else:
    guidance.append(f"Congratulations! Your loan application is eligible with a low default probab

return guidance

print("\n--- Testing User Guidance Engine ---")
# Find an example user from X_test_default that is likely to be rejected
# For testing, let's try to find a user with high default probability
rejected_user_idx = y_proba_lr.argsort()[-1] # Index of the user with highest predicted default probab
rejected_user_profile = X_test_default.iloc[rejected_user_idx]
predicted_default_prob_rejected = y_proba_lr[rejected_user_idx]

# Calculate SHAP values for this specific rejected user
shap_values_rejected_user = explainer.shap_values(rejected_user_profile)

# Generate guidance for the rejected user
guidance_output = generate_guidance(
    rejected_user_profile,
    predicted_default_prob_rejected,
    shap_values_rejected_user,
    X_test_default.columns.tolist(),
    eligibility_threshold=0.25 # Assuming this as the rejection threshold
)

print("\nGenerated Guidance for Rejected User:")
for tip in guidance_output:
    print(tip)

```

--- Testing User Guidance Engine ---

Generated Guidance for Rejected User:
 Your loan application was rejected due to a high estimated default probability of 0.29.
 Here are some personalized suggestions to improve your loan eligibility:

- Your Credit Score (302) is a major factor. Focus on paying bills on time, reducing credit card debt,

✓ WhatsApp Notification Simulation

Subtask:

Simulate the output of a WhatsApp notification for both approved and rejected loan applications. This simulation will include relevant details such as the loan amount, EMI, and any personalized guidance tips.

Reasoning: I will define the `simulate_whatsapp_notification` function as instructed, which will integrate the previously defined `estimate_loan_eligibility` and `generate_guidance` functions to create personalized WhatsApp messages for both approved and rejected loan applications. Then, I will test this function with the example eligible and rejected user profiles.

```
import textwrap

def simulate_whatsapp_notification(user_profile, credit_risk_model, shap_explainer, feature_names, eli
    """
    Simulates a WhatsApp notification for a loan application, including eligibility status and guidance.

    Args:
        user_profile (pd.Series): The user's financial profile.
        credit_risk_model: The trained machine learning model for default prediction.
        shap_explainer: The SHAP explainer for the credit risk model.
        feature_names (list): List of feature names corresponding to SHAP values.
        eligibility_threshold (float): Maximum default probability for a loan to be eligible.

    Returns:
        str: A formatted WhatsApp notification message.
    """
    # 2a. Call the estimate_loan_eligibility function
    loan_status = estimate_loan_eligibility(user_profile, credit_risk_model, eligibility_threshold)

    message = []
    message.append("🏠 *Loan Application Update* 🏠 ")
    message.append("") # For a new line

    if loan_status['is_eligible']:
        # 2b. If the loan is eligible: construct an approval message
        message.append(f"Great news, {user_profile.name if user_profile.name else 'Applicant'}!")
        message.append("Your loan application has been *APPROVED*! 🎉")
        message.append("")
        message.append("Here are your loan details:")
        message.append(f"💰 *Eligible Loan Amount:* ${loan_status['eligible_loan_amount']:.2f}")
        message.append(f"📅 *Loan Tenure:* {loan_status['loan_tenure_months']} months")
        message.append(f"💳 *Monthly EMI:* ${loan_status['monthly_emi']:.2f}")
        message.append("")
        message.append("We're excited to support your financial journey!")
        message.append("Reply 'YES' to accept or visit our portal for more options.")
    else:
        # 2c. If the loan is rejected: construct a rejection message with guidance
        message.append(f"Hello, {user_profile.name if user_profile.name else 'Applicant'}.")
        message.append("We regret to inform you that your loan application could not be *APPROVED* at")
        message.append(f"Your estimated default probability was {loan_status['default_probability']:.2f}")
        message.append("")

    # Calculate SHAP values for this specific user to generate guidance
    # Reshape user_profile to 2D array if it's a Series for explainer.shap_values
```

```

user_profile_df = pd.DataFrame([user_profile], columns=user_profile.index)
shap_values_user = shap_explainer.shap_values(user_profile_df)[0] # Get SHAP values for the si

guidance_tips = generate_guidance(
    user_profile,
    loan_status['default_probability'],
    shap_values_user,
    feature_names,
    eligibility_threshold
)

# Filter out the initial rejection message as we already have it
guidance_tips = [tip for tip in guidance_tips if "Your loan application was rejected" not in t

if guidance_tips:
    message.append("However, we'd like to help you improve your chances in the future. Here ar
    message.extend(guidance_tips)
else:
    message.append("We recommend reviewing your overall financial health to identify areas for

message.append("")
message.append("For more details, please contact our support team.")

return "\n".join(message)

print("--- Testing WhatsApp Notification for Eligible User ---")
# Use the example_user_profile from 'Loan Eligibility Estimator' subtask
# Name the Series for better personalized messages
example_user_profile.name = "Alice Smith"
eligible_notification = simulate_whatsapp_notification(
    example_user_profile,
    lr_model,
    explainer,

```